

Silencing by Design ? Copyright Enforcement, Platform Governance, and Democratic Participation in Kenya’s Digital Public Sphere A Discussion Paper on Platform Governance and Democratic Communication

Peter Ngugi

July 12, 2026

Abstract

This paper examines the intersection of copyright enforcement, platform governance, and democratic participation in Kenya’s digital public sphere. Drawing on documented cases from April 2024 and the #RejectFinanceBill2024 protests, it argues that regardless of what triggers a given mute a legitimate rights-holder match, an overzealous automated classifier, or a bad-faith report the effect of silencing political audio is a measurable loss to democratic communication. The paper deliberately stops short of alleging that any specific act of muting was intentional or politically motivated; its claim is about consequences, not motives. It situates this problem within Kenya’s constitutional guarantees of expression, access to information, media freedom, and assembly, and proposes a governance framework for platforms, regulators, and citizens that balances intellectual-property enforcement against those guarantees.

1 Introduction

In April 2024, and again during the #RejectFinanceBill2024 protests, Kenyans documented a recurring phenomenon on Facebook, TikTok, and YouTube: videos of political rallies, press briefings, and parliamentary proceedings uploaded with the picture intact but the sound stripped out. A caption would appear “audio removed due to a copyright claim” and the speech itself would vanish into silence. For a viewer scrolling past, the effect looks identical whether it was triggered by a rights-holder’s automated bot flagging a jingle playing in the background, a platform’s own moderation system misreading a crowd chant, or a third party filing a report for reasons that have nothing to do with music at all.

This paper does not claim to know which of those mechanisms was responsible in any specific case, and it deliberately avoids naming individuals or alleging

that any politician or party has been targeted. What it argues instead is narrower and, I think, more defensible: regardless of the cause, the effect of muting political audio is a measurable loss to democratic communication a loss of transparency, an increase in public suspicion, and an opening for misinformation to fill the silence. The paper treats this as a governance problem, not a conspiracy, and asks what follows from that for platforms, regulators, and citizens in Kenya’s digital public sphere.

2 Platform Governance and Automated Copyright Systems

Copyright enforcement on major platforms is now overwhelmingly automated. YouTube’s Content ID, and comparable systems on Meta’s platforms, scan uploaded audio and video against a database of registered works and act on matches without human review in the first instance. When a match is found, the standard outcomes are muting the affected audio, blocking monetisation, restricting the video’s reach, or, in more severe cases, removal following a formal notice. A copyright claim commonly results in the claimed audio segment being muted, and if the video was being monetised, that revenue is redirected to the rights holder rather than the uploader. [1] Disputing the claim requires the uploader to prove a licence or permission exists a burden that is trivial for a large broadcaster with a legal department and considerable for an independent journalist or a citizen filming with a phone.

Two structural features of these systems matter for political content specifically. First, they are content-blind in a narrow technical sense: a Content ID match is triggered by an audio fingerprint, not by an assessment of newsworthiness or public interest, so a campaign rally with a snippet of licensed music playing in the background can be muted in exactly the same way as a pirated film upload. Second, the systems are asymmetric in who can use them: rights holders and, separately, anyone able to file a take down or reporting request, hold disproportionate power over what is heard, while the burden of dispute and delay falls on the uploader. Kenyan platform governance research has already flagged this asymmetry in the adjacent context of government content-removal requests.

Government requests for removal typically cite defamation, privacy violations, impersonation, hate speech, national security, and content deemed capable of inciting unrest, and similar patterns exist elsewhere involving copyright infringement and posts perceived as threats to political stability. [2] Copyright takedown and muting sit inside this same broader ecosystem of automated and semi-automated content governance, but with even less transparency, because copyright systems are privately administered and rarely subject to public reporting in the way government requests sometimes are.

3 The #RejectFinanceBill2024 Movement and Digital Activism

The #RejectFinanceBill2024 movement emerged as a grassroots response to proposed tax increases in Kenya’s Finance Bill 2024. Citizens took to social media platforms to organise protests, share information, and document government responses. The hashtag became a national rallying point, trending domestically and drawing international attention to the unfolding legislative and political crisis.

Participants in the movement reported that videos documenting protests, parliamentary debates, and public statements were frequently muted or had their audio removed by automated content moderation systems. The affected content included:

- Videos of parliamentary proceedings where MPs debated the bill
- Citizen journalism documenting protest activities
- Press conferences by opposition leaders
- Public statements by civil society organisations
- Educational content explaining the bill’s provisions

As with the wider pattern described in Section 1, the paper does not assert that any of these individual instances was deliberately targeted. What the #RejectFinanceBill2024 experience does provide is a concentrated, dated case in which the general phenomenon of copyright-driven audio muting collided directly with a live, high-stakes episode of democratic mobilisation making it a useful lens for the governance analysis that follows.

4 Political Speech as Democratic Infrastructure

Political speech is not ordinary content; it is part of the infrastructure that lets citizens hold power to account. Kenya’s Constitution treats it this way. Article 33 is described as the bedrock of digital rights in Kenya, guaranteeing every person’s right to seek, receive, and impart information or ideas, a protection that extends to online platforms and social media. [3] That right is deliberately broad and includes narrow, specific carve-outs it does not extend to propaganda for war, incitement to violence, hate speech, or advocacy of hatred based on ethnic incitement or the grounds of discrimination in Article 27(4) [4] but ordinary political oratory, debate, and rally speech is squarely inside the zone the Constitution protects, not outside it.

The right to “receive” information, in particular, is what makes muted audio a live constitutional concern rather than a merely technical inconvenience. A citizen’s ability to hear a candidate’s own words, unmediated, is part of how they assess character, sincerity, and policy commitments judgments that are

much harder to make from a text summary or a third-party paraphrase. When the audio layer of political communication becomes unreliable, the public is pushed toward secondary sources: press coverage, partisan commentary, and social media discussion threads, each of which introduces its own selection and framing. This does not mean those secondary sources are dishonest; it means the public loses a degree of direct, low-friction access to primary political speech, and that loss is itself a democratic cost, independent of anyone’s intent.

5 Case Study: Muted Political Videos in Kenya

Kenya’s information environment has been under particular strain in recent election and protest cycles, which makes it a useful, if imperfect, setting to examine how automated moderation interacts with political content. During the 2022 general election, digital platforms were reported to be inundated with misinformation and propaganda spread by both media houses and influential social media users, with disinformation campaigns run by presidential front-runners’ own campaign teams that platforms failed to adequately curb. [5] TikTok in particular was found to have carried manipulated content, including a falsified news bulletin about a nonexistent opinion poll and fabricated newspaper covers. [5] Separately, during the 2024 Finance Bill protests, Kenya experienced a documented internet shutdown and government pressure on media houses. Citizens were unable to access critical information or verify facts during the crisis, contributing to a climate of uncertainty in which one online campaign falsely accused local media of misrepresenting the bill’s financial burden. [6] These episodes were not primarily copyright-driven, but they establish the backdrop: Kenya’s digital public sphere during politically sensitive periods is already fragile, already contested, and already prone to disinformation filling informational vacuums. Muted audio whatever its cause lands in this same fragile environment. A viewer who encounters a silenced political video during an election period is not evaluating that video in a neutral information context; they are evaluating it against a recent history of shutdowns, disputed platform moderation, and circulating manipulated content. That context is precisely why perception, discussed in Section 9, does so much independent work here, separate from whatever actually caused the mute.